

HR_Analytics_SMOTE

May 19, 2021

1 HR Analytics - Capstone Project

HR Analysis is predictive analysis to identifying the employees most likely to get promoted.

HR Process: First identify a set of employees based on recommendations/ past performance. Selected employees go through the separate training and evaluation program for each vertical. These programs are based on the required skill of each vertical. At the end of the program, based on various factors such as training performance, KPI completion (only employees with KPIs completed greater than 60% are considered) etc., employee gets promotion

For above mentioned process, the final promotions are only announced after the evaluation and this leads to delay in transition to their new roles. Hence, company needs help to identifying the eligible candidates at a particular checkpoint so that they can expedite the entire promotion cycle.

They have provided multiple attributes around Employee's past and current performance along with demographics. Now, The task is to predict whether a potential promotee at checkpoint in the test set will be promoted or not after the evaluation process.

This dataset contains 'employee_id', 'department', 'region', 'education', 'gender', 'recruitment_channel', 'no_of_trainings', 'age', 'previous_year_rating', 'length_of_service', 'KPIs_met >80%', 'awards_won?', 'avg_training_score', 'is_promoted'

The dataset has 14 features, 54808 observations.

```
[1]: # import libraries
import pandas as pd
import numpy as np
```

```
[2]: train = pd.read_csv("train_LZdllcl.csv")
test = pd.read_csv("test_2umaH9m.csv")
```

```
[3]: train.head(3)
```

```
[3]:
```

	employee_id	department	region	education	gender	\
0	65438	Sales & Marketing	region_7	Master's & above	f	
1	65141	Operations	region_22	Bachelor's	m	
2	7513	Sales & Marketing	region_19	Bachelor's	m	

	recruitment_channel	no_of_trainings	age	previous_year_rating	\
0	sourcing	1	35	5.0	

1	other	1	30	5.0
2	sourcing	1	34	3.0

	length_of_service	KPIs_met >80%	awards_won?	avg_training_score \
0	8	1	0	49
1	4	0	0	60
2	7	0	0	50

	is_promoted
0	0
1	0
2	0

```
[4]: test.head(3)
```

```
[4]:
```

	employee_id	department	region	education	gender \
0	8724	Technology	region_26	Bachelor's	m
1	74430	HR	region_4	Bachelor's	f
2	72255	Sales & Marketing	region_13	Bachelor's	m

	recruitment_channel	no_of_trainings	age	previous_year_rating \
0	sourcing	1	24	NaN
1	other	1	31	3.0
2	other	1	31	1.0

	length_of_service	KPIs_met >80%	awards_won?	avg_training_score
0	1	1	0	77
1	5	0	0	51
2	4	0	0	47

```
[5]: train.shape
```

```
[5]: (54808, 14)
```

```
[6]: test.shape
```

```
[6]: (23490, 13)
```

```
[7]: train.describe()
```

```
[7]:
```

	employee_id	no_of_trainings	age	previous_year_rating \
count	54808.000000	54808.000000	54808.000000	50684.000000
mean	39195.830627	1.253011	34.803915	3.329256
std	22586.581449	0.609264	7.660169	1.259993
min	1.000000	1.000000	20.000000	1.000000
25%	19669.750000	1.000000	29.000000	3.000000
50%	39225.500000	1.000000	33.000000	3.000000

75%	58730.500000	1.000000	39.000000	4.000000
max	78298.000000	10.000000	60.000000	5.000000

	length_of_service	KPIs_met >80%	awards_won?	avg_training_score \
count	54808.000000	54808.000000	54808.000000	54808.000000
mean	5.865512	0.351974	0.023172	63.386750
std	4.265094	0.477590	0.150450	13.371559
min	1.000000	0.000000	0.000000	39.000000
25%	3.000000	0.000000	0.000000	51.000000
50%	5.000000	0.000000	0.000000	60.000000
75%	7.000000	1.000000	0.000000	76.000000
max	37.000000	1.000000	1.000000	99.000000

	is_promoted
count	54808.000000
mean	0.085170
std	0.279137
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

```
[8]: test.describe()
```

	employee_id	no_of_trainings	age	previous_year_rating \
count	23490.000000	23490.000000	23490.000000	21678.000000
mean	39041.399149	1.254236	34.782929	3.339146
std	22640.809201	0.600910	7.679492	1.263294
min	3.000000	1.000000	20.000000	1.000000
25%	19370.250000	1.000000	29.000000	3.000000
50%	38963.500000	1.000000	33.000000	3.000000
75%	58690.000000	1.000000	39.000000	4.000000
max	78295.000000	9.000000	60.000000	5.000000

	length_of_service	KPIs_met >80%	awards_won?	avg_training_score
count	23490.000000	23490.000000	23490.000000	23490.000000
mean	5.810387	0.358834	0.022776	63.263133
std	4.207917	0.479668	0.149191	13.411750
min	1.000000	0.000000	0.000000	39.000000
25%	3.000000	0.000000	0.000000	51.000000
50%	5.000000	0.000000	0.000000	60.000000
75%	7.000000	1.000000	0.000000	76.000000
max	34.000000	1.000000	1.000000	99.000000

```
[9]: train.isnull().sum()
```

```
[9]: employee_id      0
      department      0
      region          0
      education      2409
      gender          0
      recruitment_channel  0
      no_of_trainings  0
      age            0
      previous_year_rating 4124
      length_of_service  0
      KPIs_met >80%    0
      awards_won?      0
      avg_training_score  0
      is_promoted      0
      dtype: int64
```

```
[10]: test.isnull().sum()
```

```
[10]: employee_id      0
      department      0
      region          0
      education      1034
      gender          0
      recruitment_channel  0
      no_of_trainings  0
      age            0
      previous_year_rating 1812
      length_of_service  0
      KPIs_met >80%    0
      awards_won?      0
      avg_training_score  0
      dtype: int64
```

```
[11]: train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 54808 entries, 0 to 54807
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   employee_id           54808 non-null  int64
1   department            54808 non-null  object
2   region               54808 non-null  object
3   education             52399 non-null  object
4   gender               54808 non-null  object
5   recruitment_channel   54808 non-null  object
6   no_of_trainings      54808 non-null  int64
```

```

7   age                54808 non-null  int64
8   previous_year_rating  50684 non-null  float64
9   length_of_service    54808 non-null  int64
10  KPIs_met >80%        54808 non-null  int64
11  awards_won?          54808 non-null  int64
12  avg_training_score    54808 non-null  int64
13  is_promoted           54808 non-null  int64
dtypes: float64(1), int64(8), object(5)
memory usage: 5.9+ MB

```

```
[12]: train.education.value_counts()
```

```

[12]: Bachelor's          36669
      Master's & above    14925
      Below Secondary      805
      Name: education, dtype: int64

```

```
[14]: train['education'].fillna("Bachelor's", inplace=True)
      test['education'].fillna("Bachelor's", inplace=True)
```

```
[15]: train['previous_year_rating'].skew()
```

```
[15]: -0.3106378431385327
```

```
[16]: train['previous_year_rating'].value_counts()
```

```

[16]: 3.0    18618
      5.0    11741
      4.0     9877
      1.0     6223
      2.0     4225
      Name: previous_year_rating, dtype: int64

```

```

[17]: train['previous_year_rating'].fillna(train['previous_year_rating'].median(),
      ↪inplace=True)
      test['previous_year_rating'].fillna(test['previous_year_rating'].median(),
      ↪inplace=True)

```

```
[18]: train.isnull().sum()
```

```

[18]: employee_id      0
      department      0
      region          0
      education        0
      gender           0
      recruitment_channel 0
      no_of_trainings  0

```

```

age                0
previous_year_rating 0
length_of_service  0
KPIs_met >80%      0
awards_won?        0
avg_training_score  0
is_promoted         0
dtype: int64

```

```
[19]: test.isnull().sum()
```

```

[19]: employee_id      0
      department      0
      region          0
      education        0
      gender           0
      recruitment_channel 0
      no_of_trainings  0
      age              0
      previous_year_rating 0
      length_of_service 0
      KPIs_met >80%    0
      awards_won?      0
      avg_training_score 0
      dtype: int64

```

```
[20]: train.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 54808 entries, 0 to 54807
Data columns (total 14 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   employee_id           54808 non-null  int64
 1   department            54808 non-null  object
 2   region                54808 non-null  object
 3   education              54808 non-null  object
 4   gender                54808 non-null  object
 5   recruitment_channel    54808 non-null  object
 6   no_of_trainings        54808 non-null  int64
 7   age                   54808 non-null  int64
 8   previous_year_rating   54808 non-null  float64
 9   length_of_service      54808 non-null  int64
10  KPIs_met >80%          54808 non-null  int64
11  awards_won?            54808 non-null  int64
12  avg_training_score      54808 non-null  int64
13  is_promoted            54808 non-null  int64

```

```
dtypes: float64(1), int64(8), object(5)
memory usage: 5.9+ MB
```

```
[21]: from sklearn import preprocessing
      le = preprocessing.LabelEncoder()
```

```
[22]: train['department'] = le.fit_transform(train['department'])
      test['department'] = le.fit_transform(test['department'])
      train['gender'] = le.fit_transform(train['gender'])
      test['gender'] = le.fit_transform(test['gender'])
      train['education'] = le.fit_transform(train['education'])
      test['education'] = le.fit_transform(test['education'])
      train['recruitment_channel'] = le.fit_transform(train['recruitment_channel'])
      test['recruitment_channel'] = le.fit_transform(test['recruitment_channel'])
```

```
[23]: train.drop(labels='employee_id',axis=1,inplace=True)
      train.drop(labels='region',axis=1,inplace=True)
      test.drop(labels='region',axis=1,inplace=True)
```

```
[24]: train.head(3)
```

```
[24]:
```

	department	education	gender	recruitment_channel	no_of_trainings	age	\
0	7	2	0	2	1	35	
1	4	0	1	0	1	30	
2	7	0	1	2	1	34	

	previous_year_rating	length_of_service	KPIs_met >80%	awards_won?	\
0	5.0	8	1	0	
1	5.0	4	0	0	
2	3.0	7	0	0	

	avg_training_score	is_promoted
0	49	0
1	60	0
2	50	0

```
[25]: train.columns
```

```
[25]: Index(['department', 'education', 'gender', 'recruitment_channel',
        'no_of_trainings', 'age', 'previous_year_rating', 'length_of_service',
        'KPIs_met >80%', 'awards_won?', 'avg_training_score', 'is_promoted'],
        dtype='object')
```

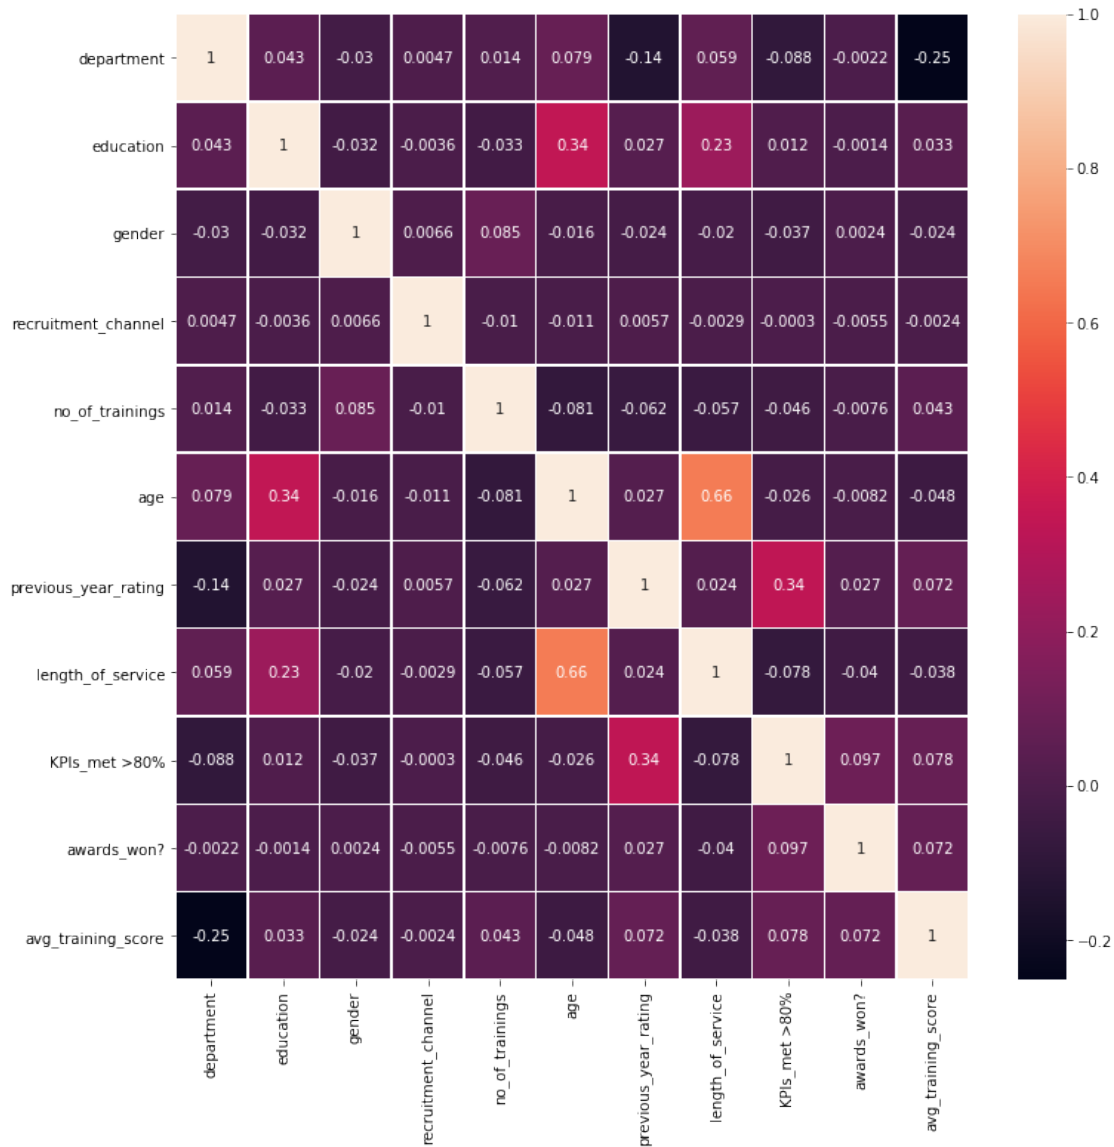
```
[26]: rel_feat = ['department', 'education', 'gender', 'recruitment_channel',
        'no_of_trainings', 'age', 'previous_year_rating', 'length_of_service',
        'KPIs_met >80%', 'awards_won?', 'avg_training_score']
```

```
[27]: rel_feat_corr = train.corr()['is_promoted'][['department', 'education', 'gender', 'recruitment_channel', 'no_of_trainings', 'age', 'previous_year_rating', 'length_of_service', 'KPIs_met >80%', 'awards_won?', 'avg_training_score']]
```

```
[28]: rel_feat_corr
```

```
[28]: department          0.000130
education                0.029257
gender                  -0.011109
recruitment_channel      0.002229
no_of_trainings         -0.024896
age                    -0.017166
previous_year_rating     0.153230
length_of_service       -0.010670
KPIs_met >80%           0.221582
awards_won?             0.195871
avg_training_score       0.181147
Name: is_promoted, dtype: float64
```

```
[30]: import matplotlib.pyplot as plt
      %matplotlib inline
      import seaborn as sns
      plt.figure(figsize = (12,12))
      sns.heatmap(train[rel_feat].corr(),annot = True, linewidths = 0.5);
```

```
[31]: X= train[rel_feat]
      y= train['is_promoted']
```

```
[32]: from imblearn.over_sampling import SMOTE
      from sklearn.model_selection import train_test_split
      X_train, X_test, y_train, y_test = train_test_split(X, y, random_state = 1,
      ↳stratify=y)
```

```
[33]: y_train.value_counts()
```

```
[33]: 0    37605
      1     3501
```

Name: is_promoted, dtype: int64

```
[34]: smt = SMOTE()  
X_train, y_train = smt.fit_sample(X_train, y_train)  
np.bincount(y_train)
```

```
[34]: array([37605, 37605], dtype=int64)
```

2 Random Forest

```
[35]: #Import Random Forest Model  
from sklearn.ensemble import RandomForestClassifier  
from sklearn.feature_selection import SelectFromModel  
  
parameters = {'bootstrap': False,  
              'min_samples_leaf': 3,  
              'n_estimators': 500,  
              'min_samples_split': 10,  
              'max_features': 'sqrt',  
              'max_depth': 10,  
              }  
  
#Create a random forest classifier, 100 trees  
clf_rf=RandomForestClassifier(**parameters)  
  
#Train the model using the training sets  
clf_rf.fit(X_train,y_train)  
  
rf_pred=clf_rf.predict(X_test).astype(int)
```

```
[36]: from sklearn.metrics import classification_report, confusion_matrix,  
      ↪ accuracy_score, recall_score  
  
print(confusion_matrix(y_test,rf_pred))  
print(classification_report(y_test,rf_pred))  
print("Accuracy:",accuracy_score(y_test, rf_pred))
```

```
[[9990 2545]  
 [ 360  807]]
```

	precision	recall	f1-score	support
0	0.97	0.80	0.87	12535
1	0.24	0.69	0.36	1167
accuracy			0.79	13702
macro avg	0.60	0.74	0.62	13702

weighted avg 0.90 0.79 0.83 13702

Accuracy: 0.7879871551598306

```
[37]: recall_score(y_test, rf_pred)
```

```
[37]: 0.6915167095115681
```

```
[38]: rf_pred = clf_rf.predict(test[rel_feat]).astype(int)
sub = pd.DataFrame()
sub['employee_id'] = test['employee_id']
sub['is_promoted'] = rf_pred
sub[['employee_id', 'is_promoted']].to_csv('submission_rf.csv', index=False)
```

```
[39]: sub.head()
```

```
[39]:
```

	employee_id	is_promoted
0	8724	1
1	74430	0
2	72255	0
3	38562	0
4	64486	0

```
[40]: sub.shape
```

```
[40]: (23490, 2)
```

```
[41]: sub.is_promoted.value_counts()
```

```
[41]: 0    17774
1     5716
Name: is_promoted, dtype: int64
```

3 XGBoost

```
[42]: from xgboost import plot_importance
```

```
[43]: # XGB Classifier
from xgboost import XGBClassifier

clf_xgb = XGBClassifier(n_estimators=200,
max_depth=10,
min_child_weight=5,
gamma=0,
subsample=0.5,
#colsample_bytree=0.3,
objective= 'binary:logistic',
```

```

nthread=5,
scale_pos_weight=13,
reg_lambda=5,
alpha=5,
base_score=0.15,
#seed=1029,
random_state=45)

clf_xgb.fit(X_train, y_train)

```

```

[43]: XGBClassifier(alpha=5, base_score=0.15, booster='gbtree', colsample_bylevel=1,
        colsample_bynode=1, colsample_bytree=1, gamma=0,
        learning_rate=0.1, max_delta_step=0, max_depth=10,
        min_child_weight=5, missing=None, n_estimators=200, n_jobs=1,
        nthread=5, objective='binary:logistic', random_state=45,
        reg_alpha=0, reg_lambda=5, scale_pos_weight=13, seed=None,
        silent=None, subsample=0.5, verbosity=1)

```

```

[44]: # Predicting the Test set results
xg_pred = clf_xgb.predict(X_test).astype(int)

```

```

[45]: # evaluate predictions
print(confusion_matrix(y_test,xg_pred))
print(classification_report(y_test,xg_pred))
print("Accuracy:",accuracy_score(y_test, xg_pred))

```

```

[[9484 3051]
 [ 222  945]]

```

		precision	recall	f1-score	support
	0	0.98	0.76	0.85	12535
	1	0.24	0.81	0.37	1167
	accuracy			0.76	13702
	macro avg	0.61	0.78	0.61	13702
	weighted avg	0.91	0.76	0.81	13702

Accuracy: 0.7611297620785287

```

[46]: recall_score(y_test, xg_pred)

```

```

[46]: 0.8097686375321337

```

```

[47]: xgb_pred = clf_xgb.predict(test[rel_feat]).astype(int)
sub = pd.DataFrame()
sub['employee_id'] = test['employee_id']
sub['is_promoted'] = xgb_pred

```

```
sub[['employee_id','is_promoted']].to_csv('submission_xgb.csv',index=False)
```

```
[48]: sub.is_promoted.value_counts()
```

```
[48]: 0    16706
      1     6784
      Name: is_promoted, dtype: int64
```

4 GBM

```
[49]: from sklearn import ensemble
      gbm = ensemble.GradientBoostingClassifier(n_estimators = 500, max_depth = 10,
      ↪min_samples_split = 10, learning_rate = 0.1)
      gbm.fit(X_train,y_train)
```

```
[49]: GradientBoostingClassifier(ccp_alpha=0.0, criterion='friedman_mse', init=None,
      learning_rate=0.1, loss='deviance', max_depth=10,
      max_features=None, max_leaf_nodes=None,
      min_impurity_decrease=0.0, min_impurity_split=None,
      min_samples_leaf=1, min_samples_split=10,
      min_weight_fraction_leaf=0.0, n_estimators=500,
      n_iter_no_change=None, presort='deprecated',
      random_state=None, subsample=1.0, tol=0.0001,
      validation_fraction=0.1, verbose=0,
      warm_start=False)
```

```
[50]: gbm_pred = gbm.predict(X_test)
```

```
[51]: # evaluate predictions
      print(confusion_matrix(y_test,gbm_pred))
      print(classification_report(y_test,gbm_pred))
      print("Accuracy:",accuracy_score(y_test, gbm_pred))
```

```
[[11919   616]
 [   667   500]]
              precision    recall  f1-score   support

         0       0.95      0.95      0.95     12535
         1       0.45      0.43      0.44       1167

 accuracy                   0.91     13702
 macro avg              0.70      0.69      0.69     13702
 weighted avg           0.90      0.91      0.91     13702
```

```
Accuracy: 0.9063640344475259
```

```
[52]: recall_score(y_test, gbm_pred)
```

```
[52]: 0.4284490145672665
```

```
[53]: gbm_pred = gbm.predict(test[rel_feat]).astype(int)
sub = pd.DataFrame()
sub['employee_id'] = test['employee_id']
sub['is_promoted'] = gbm_pred
sub[['employee_id', 'is_promoted']].to_csv('submission_gbm.csv', index=False)
```

```
[54]: sub.is_promoted.value_counts()
```

```
[54]: 0    21514
1     1976
Name: is_promoted, dtype: int64
```